Cloud Computing

COSC 2639

Lecture 1

Introduction to Cloud Computing and System Architectures

Course Overview

- Offering Coordinator and Tutor: **Prabath Abeysekara**
- Email: prabath.abeysekara@rmit.edu.au
- Tutor: **Tim Lo**
- Email: tim.lo@rmit.edu.au
- Course code: **COSC2639/2697**
- 1 hr course webinar per week
- 1 hr class webinar per week

Pre-requisites, Assumed knowledge and Capabilities

- Good knowledge in programming (This is a MUST!)
- Good knowledge in Java/PHP/Python or any other programming language supported by GCP and AWS
- Basic knowledge in Net-centric Communication
- Basic knowledge in NodeJS
- Basic knowledge in Data Analytics

Important:

- Be able to access all Google (Cloud) Services (otherwise you CANNOT do Assessment 1 (30%)

Assessments

Assessment 01: Timed Practical Google Cloud Application Implementation (30%)

- ❑ Individual
- ❑ Specification released in early Week 2
- ❑ Due in Week 5
- ❑ Demonstration needed for assessment (Week 6)

Assessment 02: Timed Practical AWS Cloud Database Application Development (20%)

- ❑ Individual
- ❑ Specification released in early Week 6
- ❑ Due in Week 9
- ❑ Demonstration needed for assessment (Week 10)

Assessment 03: Timed Practical AWS Cloud System Development (50%)

- ❑ Individual
- ❑ Specification released in early Week 4
- ❑ Due in Week 12
- ❑ Demonstration needed for assessment (Week 13)

Use of Generative AI tools

- GenAI tools (e.g. ChatGPT) have become a part of everyone's learning journey, nowadays.
- Encourage using them as "assistive technologies" to make you more productive.
- Discourage using them to do the "assignments" for you! ☺

1 Cloud computing use cases

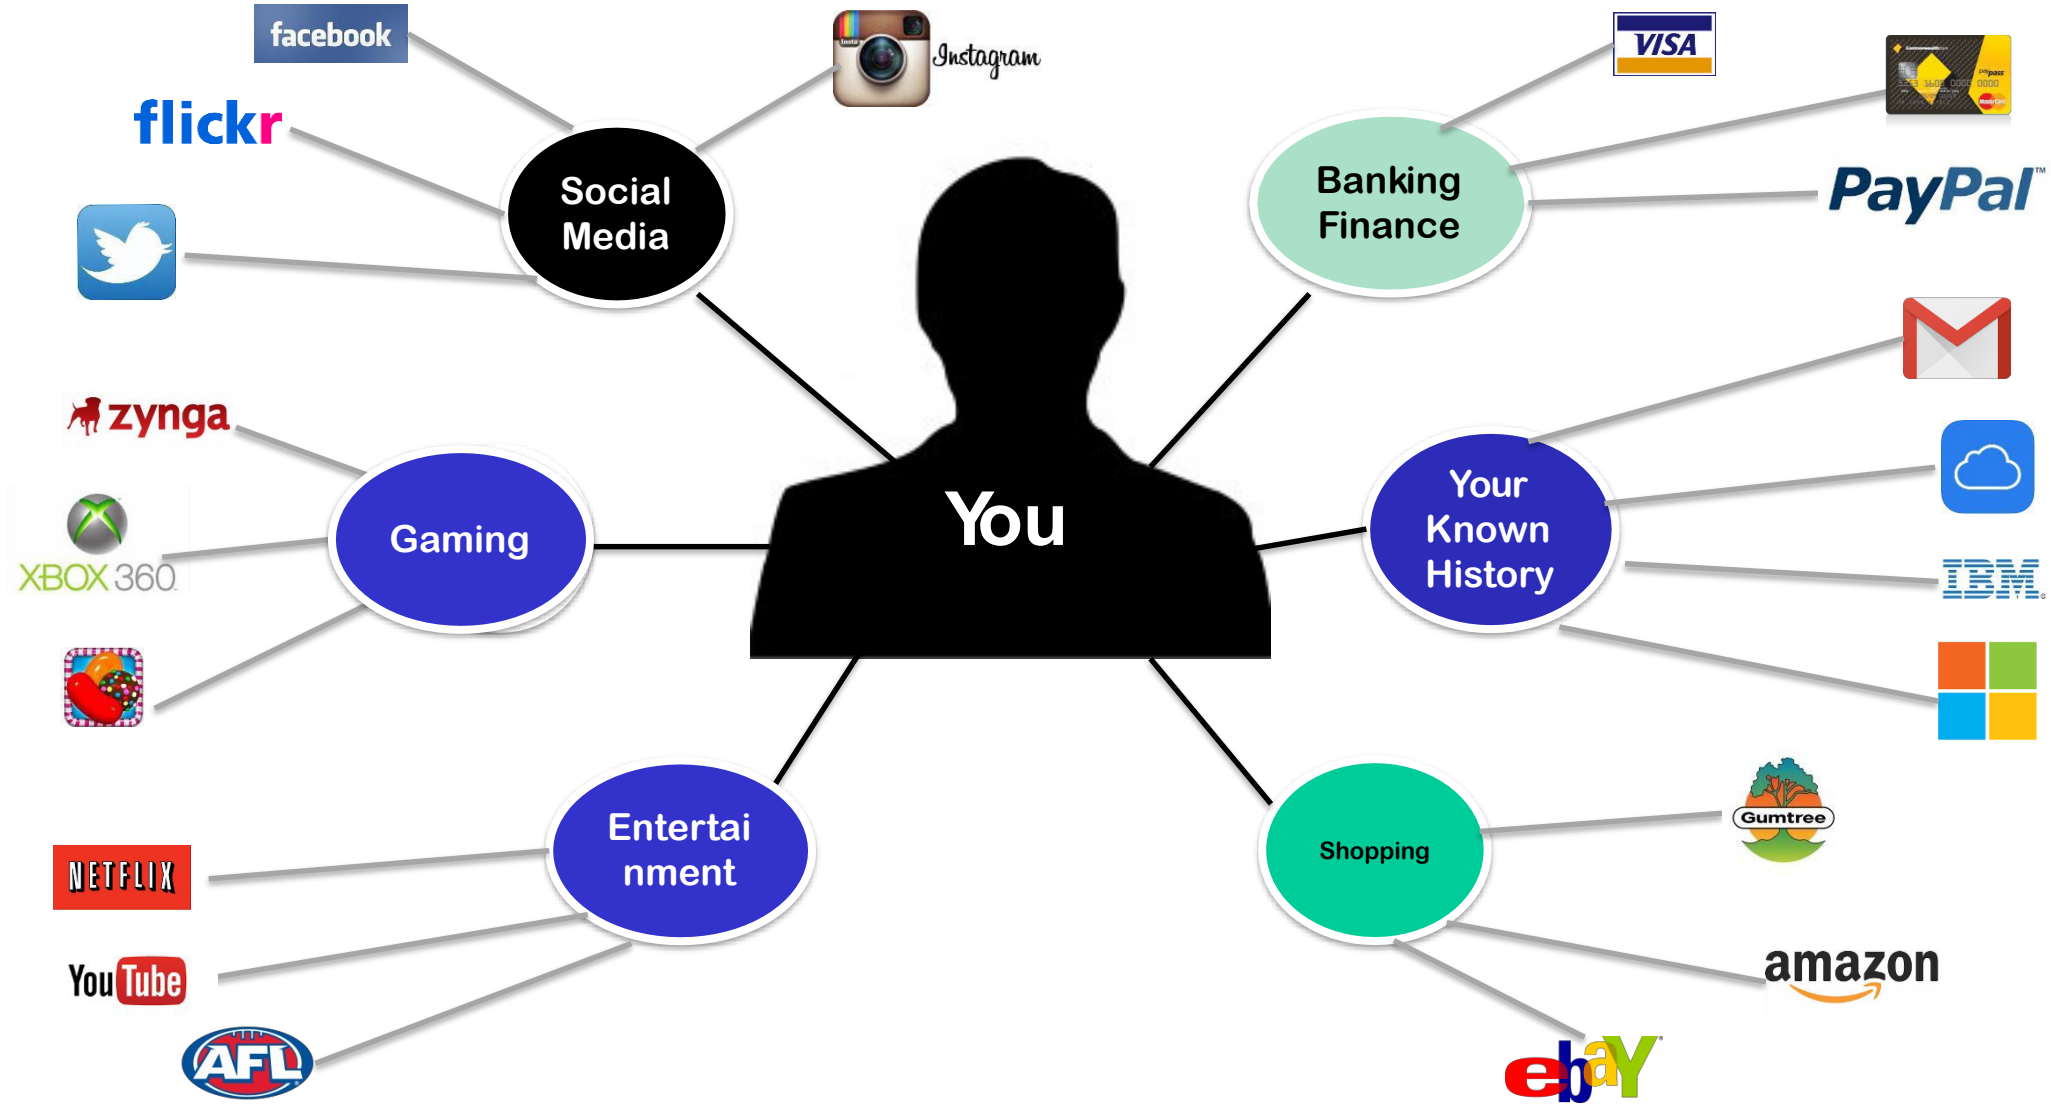

Use Case – Cloud accelerates AI: Self-driving car

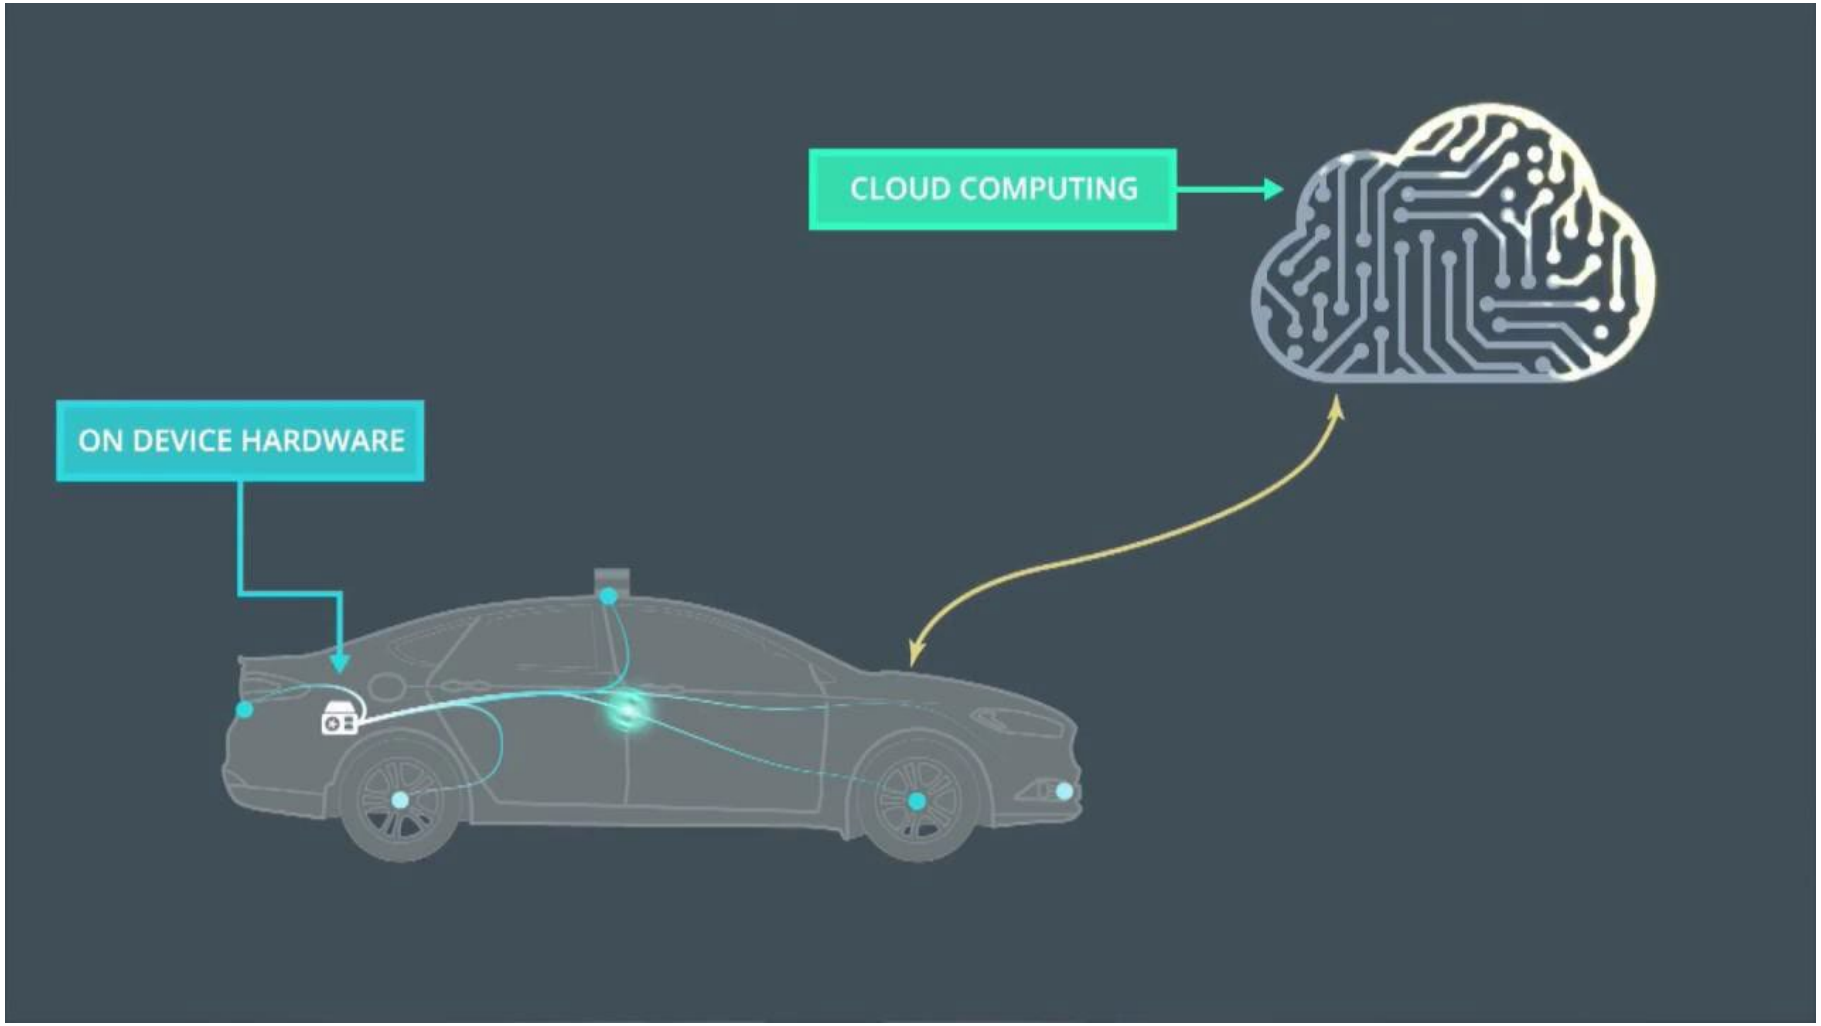

Use Case – Cloud accelerates AI: Smart city

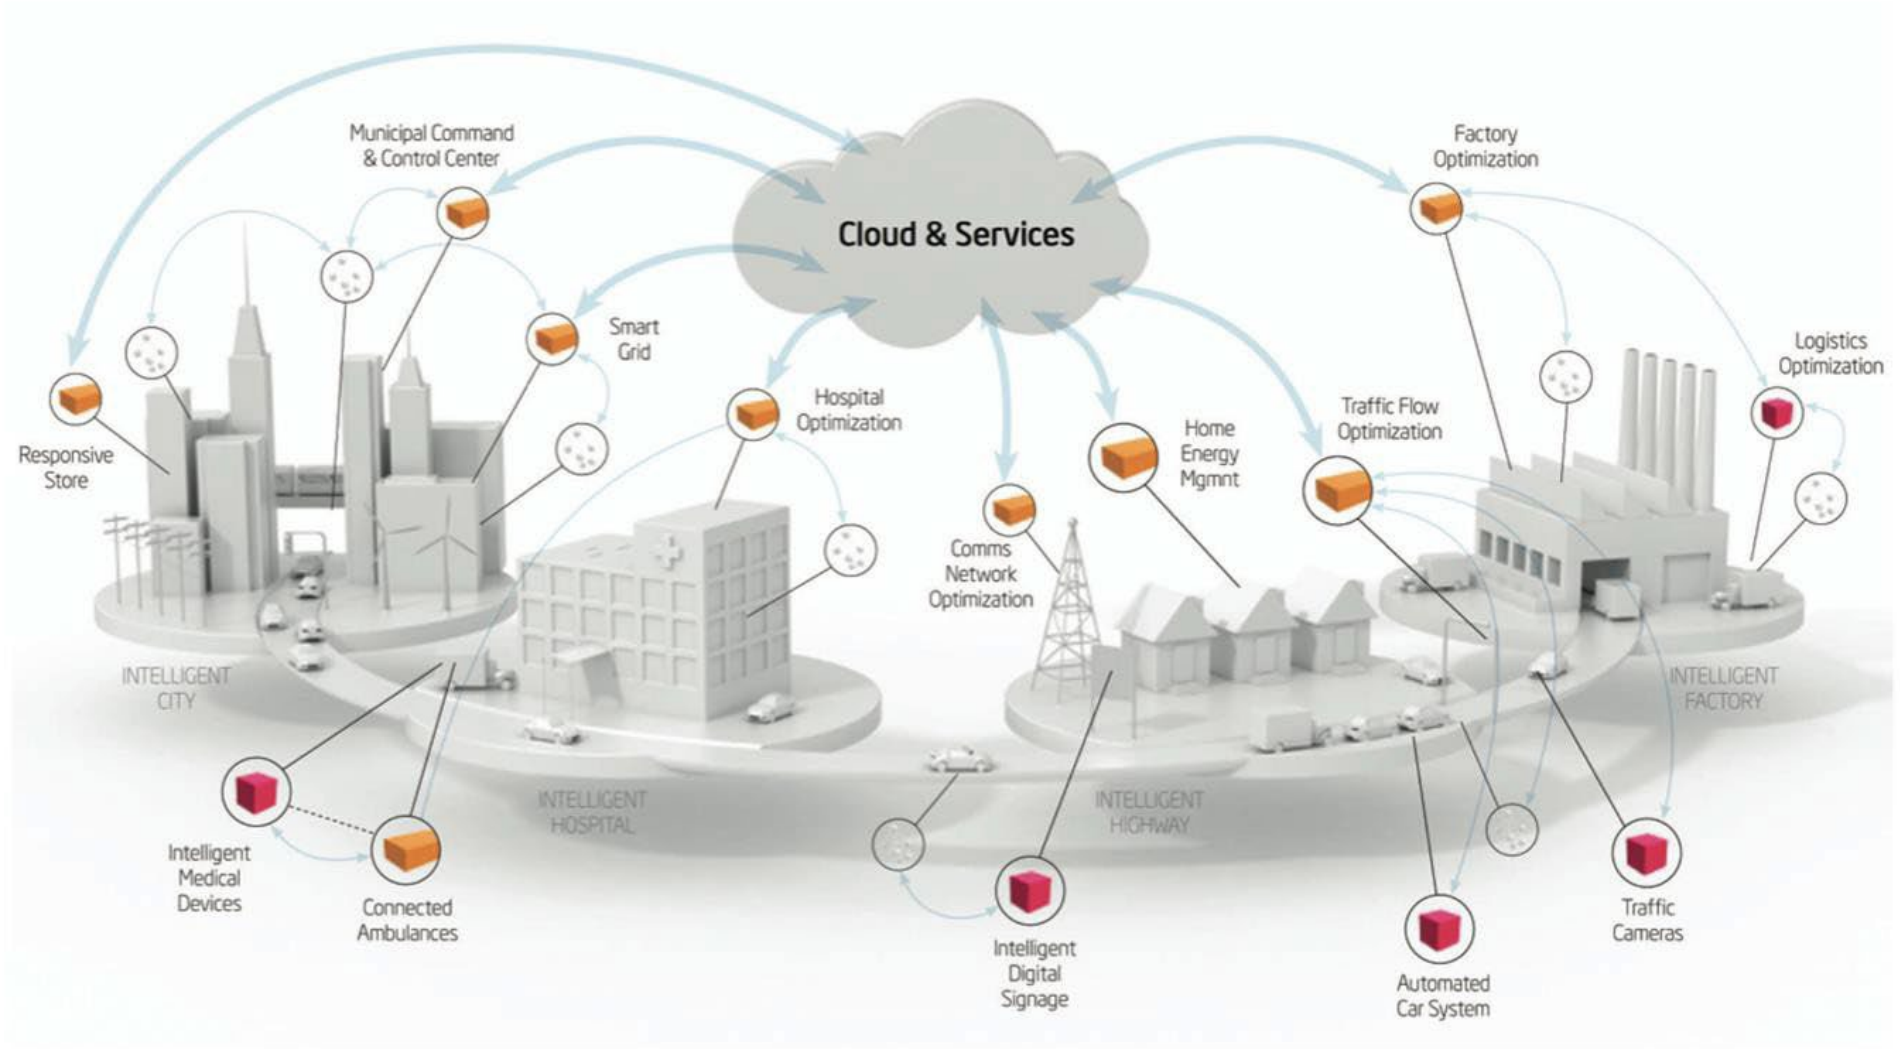

2 Cloud computing evolution

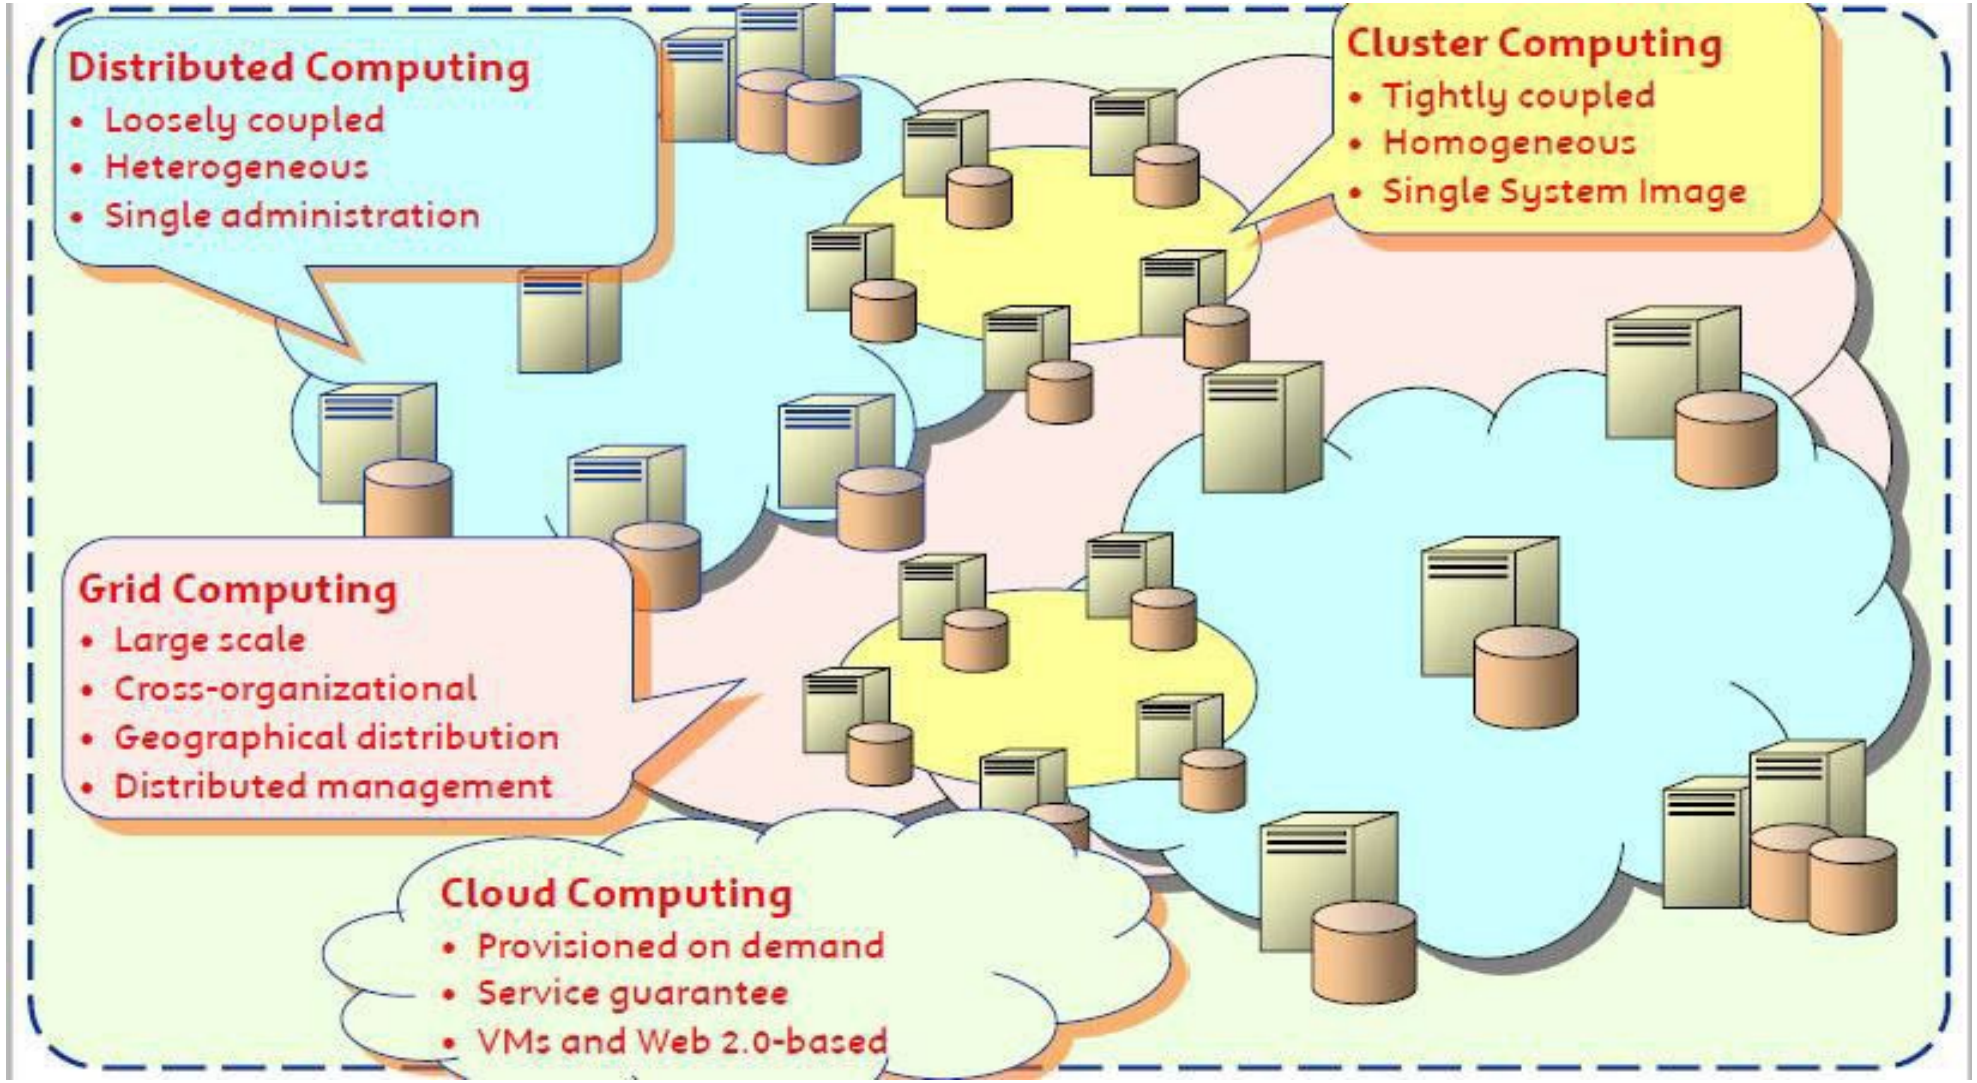

2.1 Centralized computing

- All computer resources are centralized in one physical system.
- All resources (processors, memory, and storage) are tightly coupled within one integrated OS.
- Example: A single printer in a network of computers.

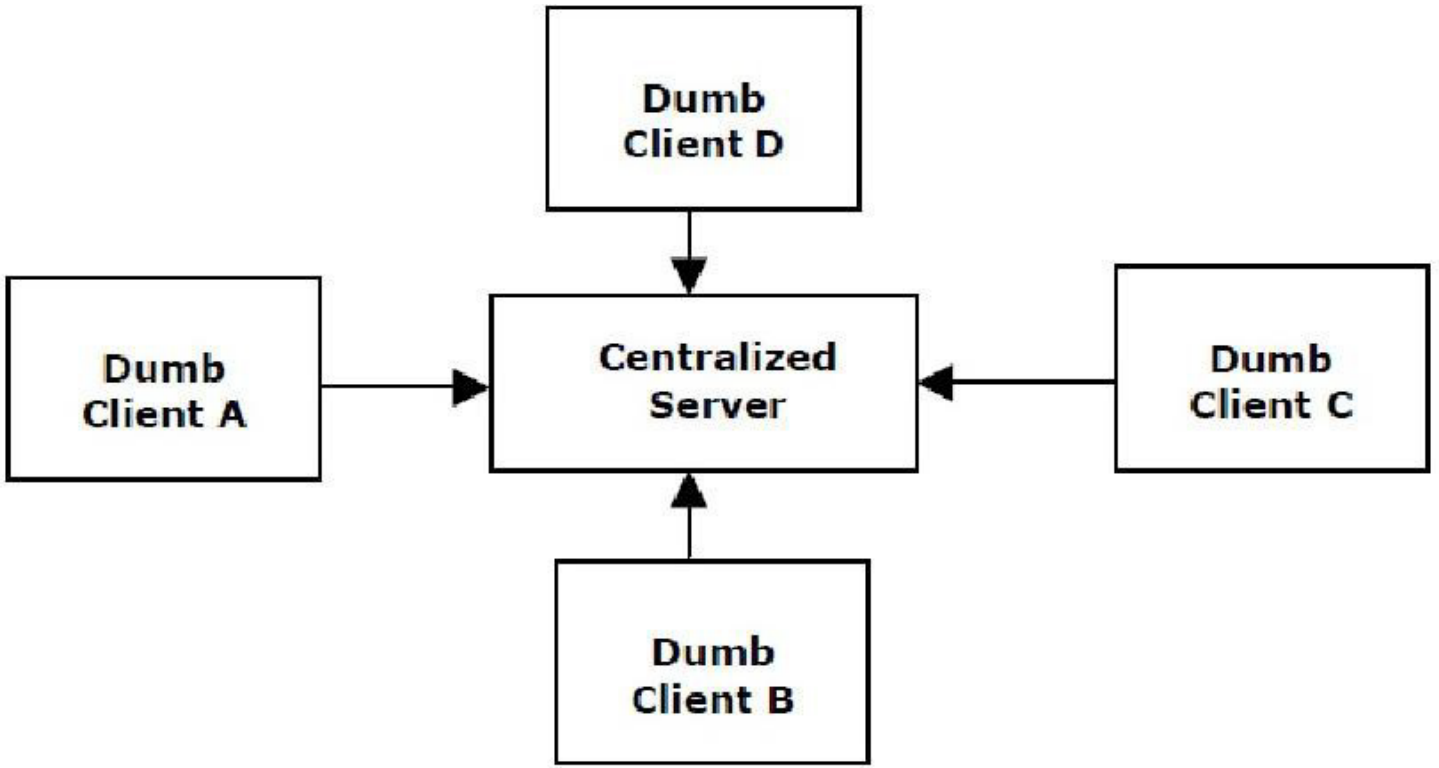

2.2 Parallel computing

- Multiple processes can run in parallel using a shared memory
- all processors are tightly coupled with centralized memory.
- Inter-processor communication is accomplished through shared memory or via message passing.
- A computer system capable of parallel computing (has multiple processors) is commonly known as a parallel computer.

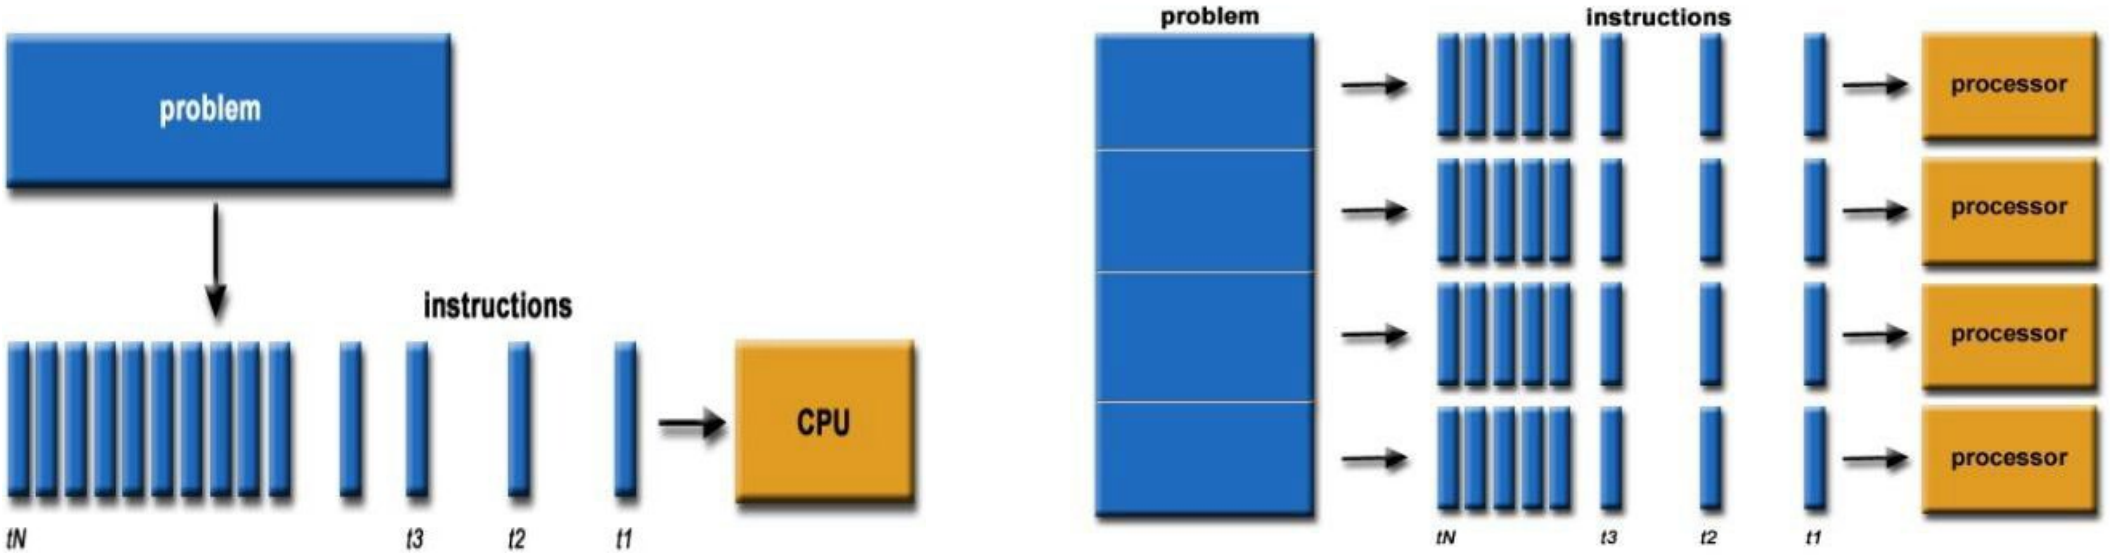

2.3 Cluster Computing

- Cluster is tightly-coupled
- Resources are homogenous or identical
- Work together for a specific job
- Each node performs the same task
- Clustering explores massive parallelism at the job level and achieves high availability (HA) through stand-alone operations.

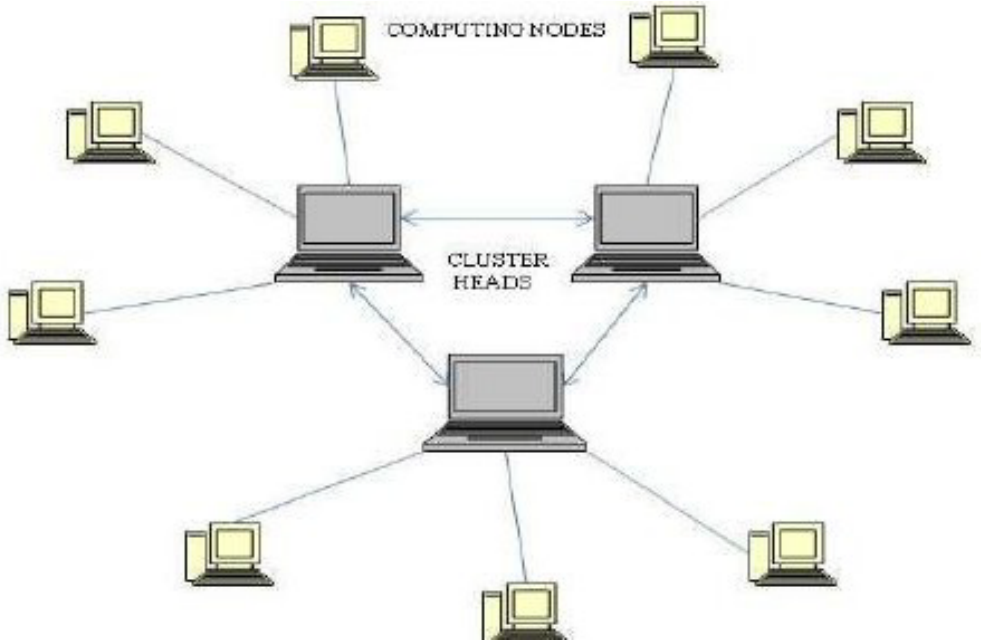

Evolution of cluster computing

2.4 Distributed computing

- Consists of multiple autonomous computers.
- Each computer has its own private memory.
- Communicate through a computer network.
- Information exchange in a distributed system is accomplished through message passing.
- Originated from distributed system is – Distributed programming

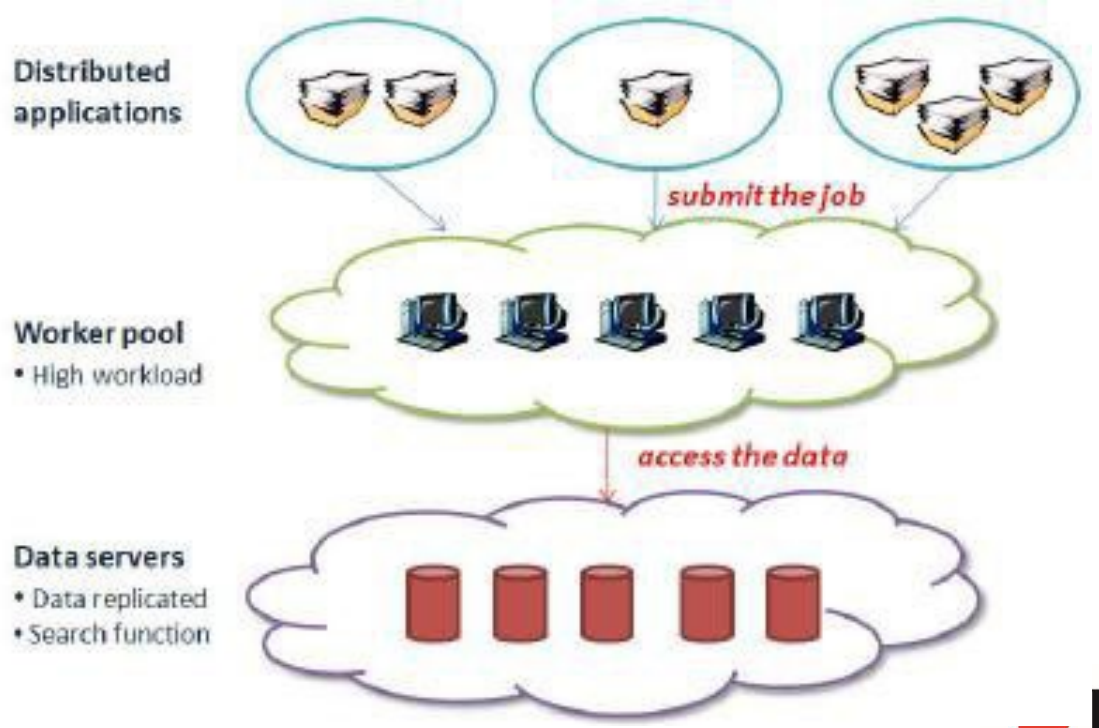

2.5 Grid computing

- Grid Computing
  - A variation of distributed computing
  - Operates at a much larger scale
  - Cross-organizational
  - Geographically distributed
  - Distributed management

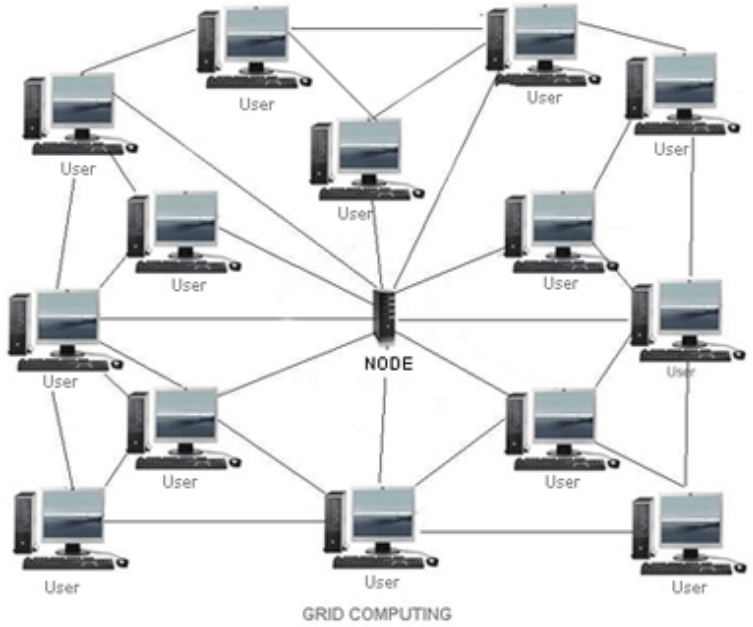

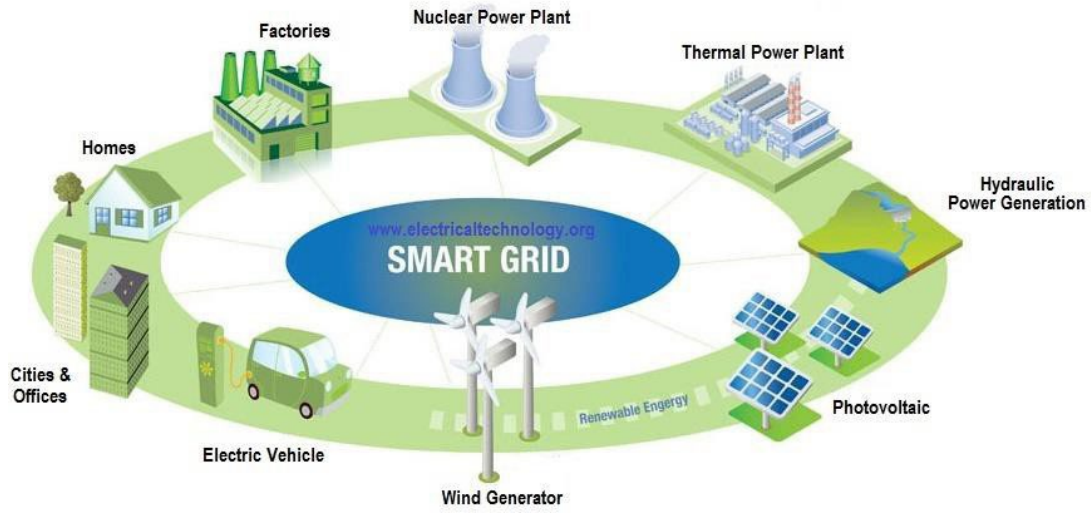

Grid Computing

- Grid systems are classified in two categories: computational or data grids and P2P grids.

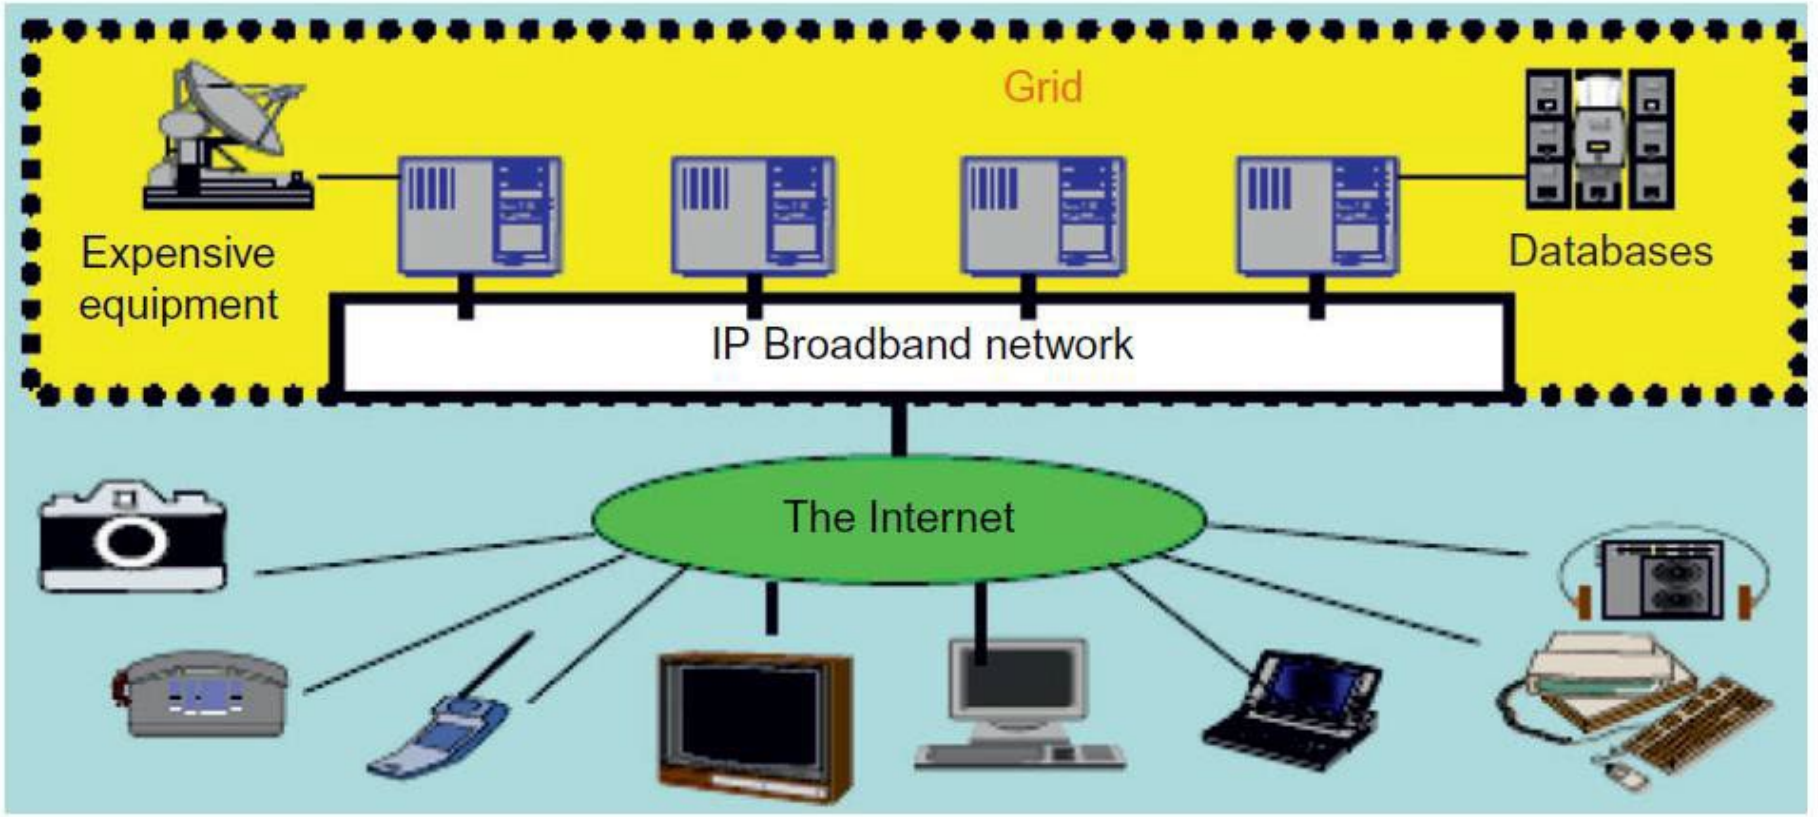

2.6 Cloud Computing

- It can be centralized or distributed. How?
  - Cloud can be built with physical or virtualized resources over large data centers that are centralized or distributed.
- Applies parallel or distributed computing or combination of both
- Provisioned on-demand
- Service guarantees
- Three layer-ed services

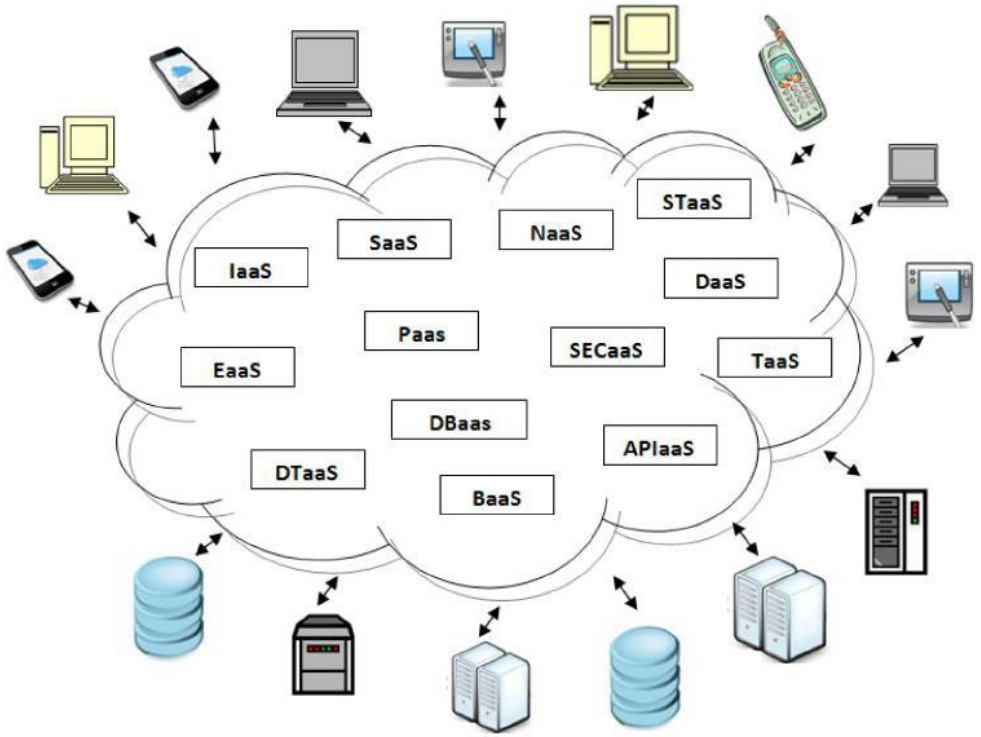

2.7 Mobile Computing

- The computing environment is mobile and moves with users.
- Example?

2.8 Ubiquitous computing

- Wireless communication between the components
- Refers to computing with pervasive devices at any place and time using wired or wireless communication.
- Also known as Pervasive computing.
- Example?
  - Smart Traffic Light
  - Home automation system
  - Fitbit!

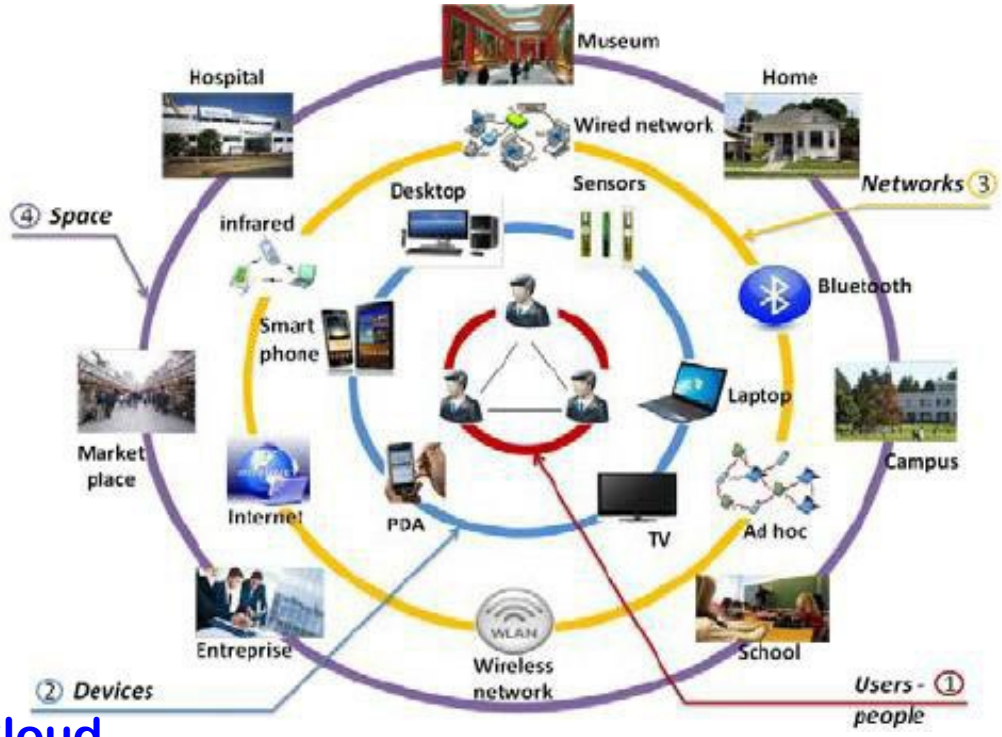

Bing Huang, Athman Bouguettaya, Hai Dong: Enabling Edge Cloud Intelligence for Activity Learning in Smart Home (2020)

2.9 Edge computing

- Local processing before reaching cloud
- Complements the concept of IoT
- Data from IoT devices are analyzed at the edge before reaching the cloud.
- Sensors, controllers, and other connected devices collect and analyze IoT data themselves, or transmit it to a near by computing computing device (server, laptop, Raspberry Pi).

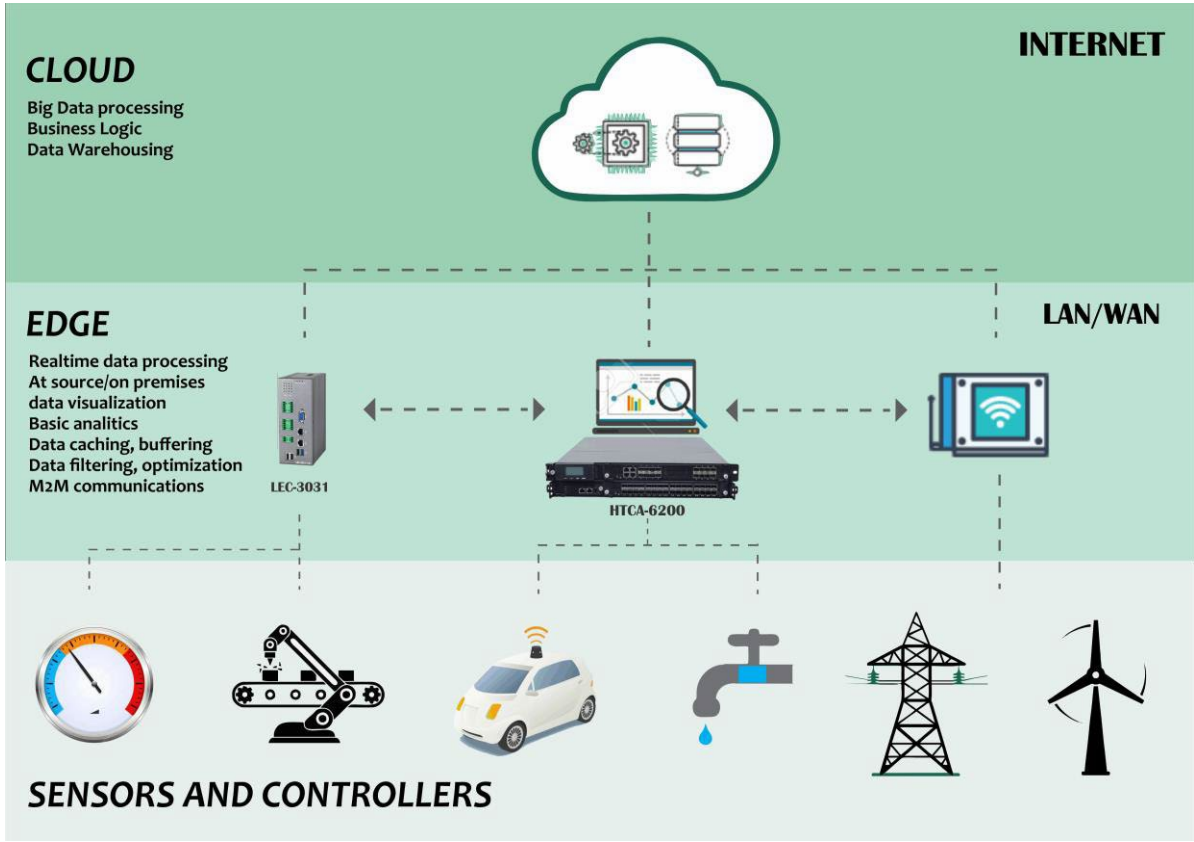

Abeysekara, P., Dong, H., Qin, A. K.: Distributed Machine Learning for Predictive Analytics in Mobile Edge Computing Based IoT Environments. The 2020 International Joint Conference on Neural Networks (IJCNN 2020) (July 2020)

2.10 Fog computing

- A small cloud before data go the actual cloud
- Enterprise network
- Provide computing, storage and networking services for data processing.

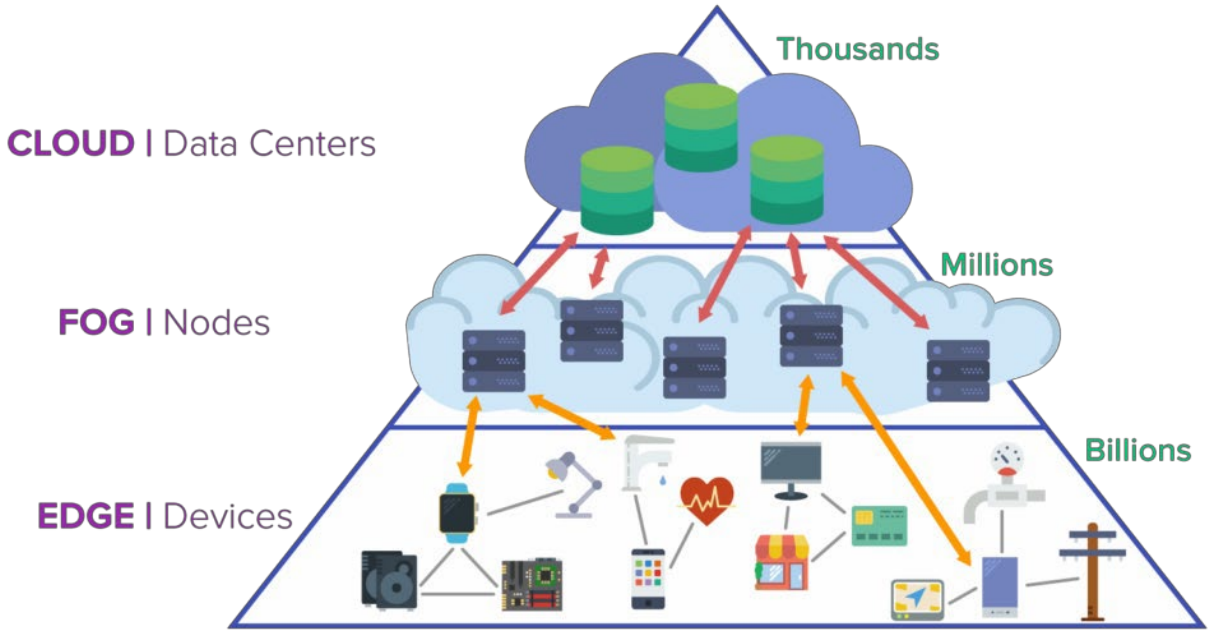

3 Essential Cloud Characteristics

- Large data centers
- Virtualization
- Pay-as-you go (reduced cost)
- Scalable and on-demand services
- Broad network access
- Resource pooling
  - Location independence
- Rapid elasticity
- Measured service

3.1 Large Data Centers

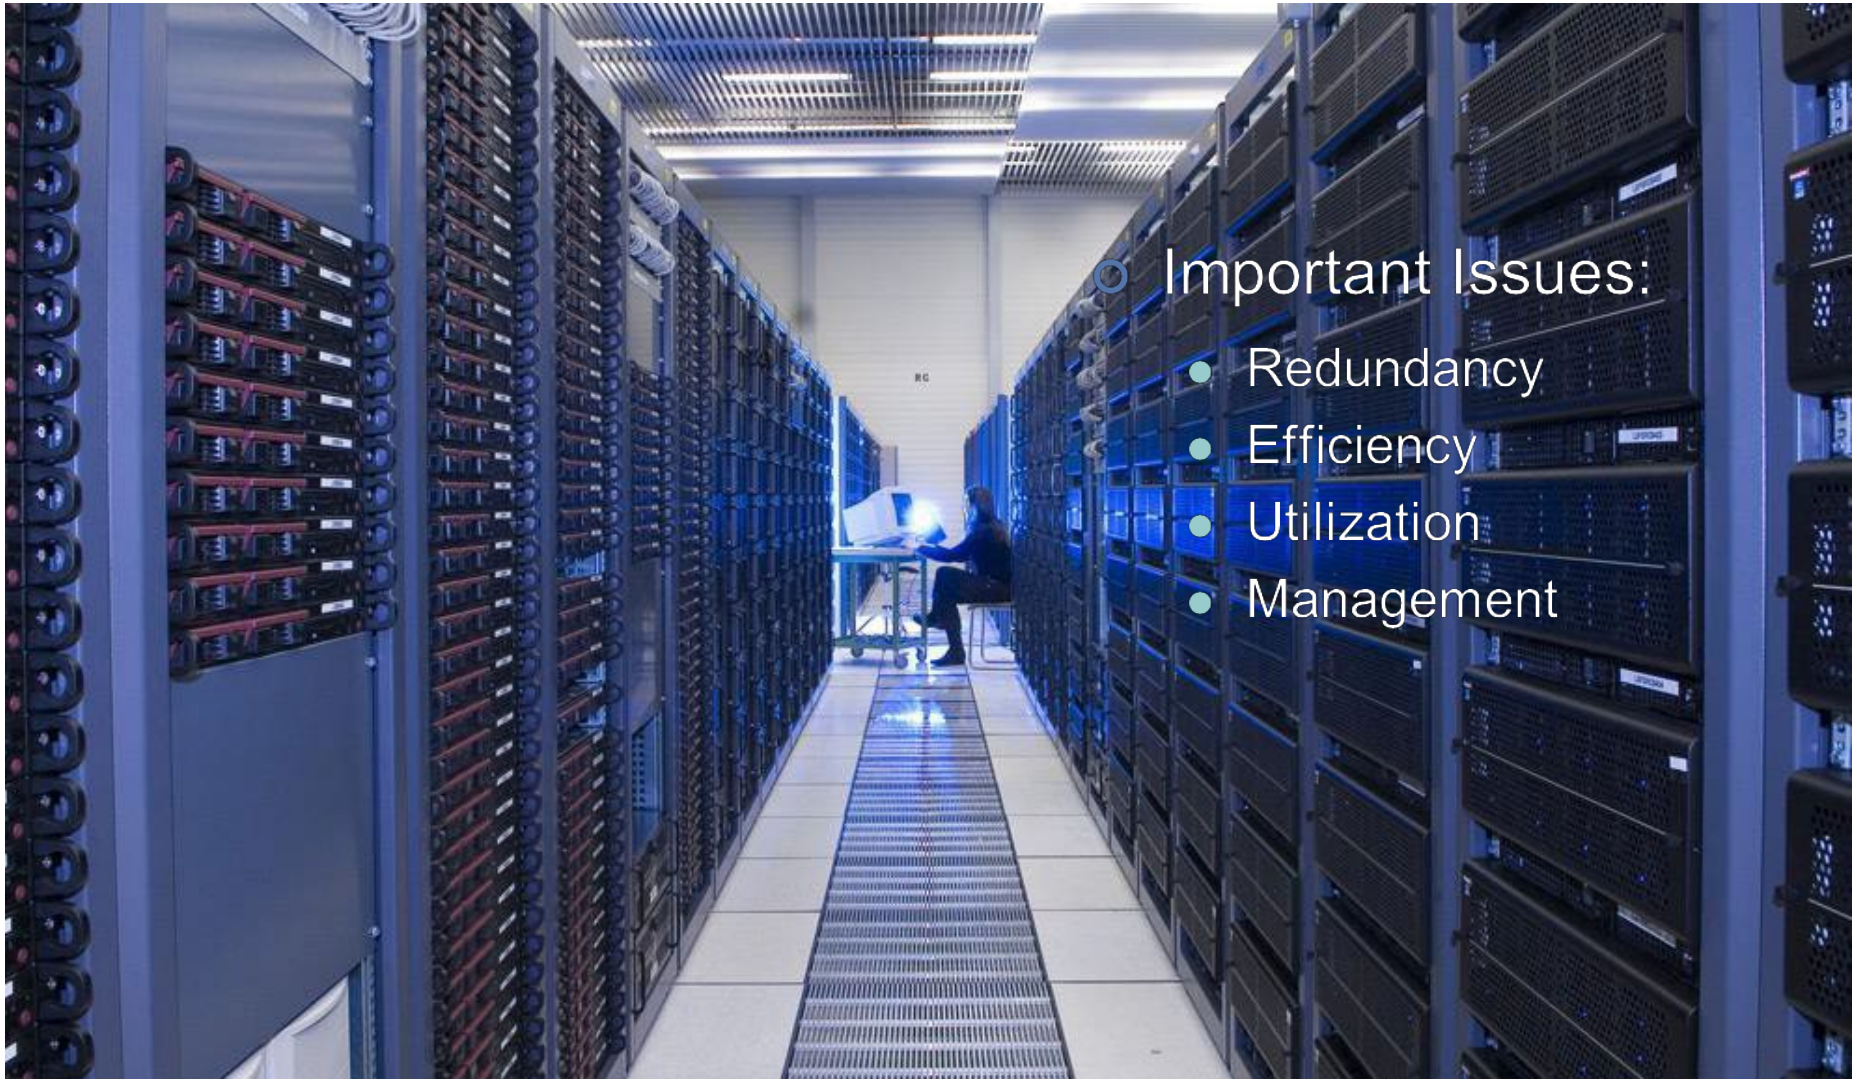

Large Data Centers

- Data centers have grown rapidly in recent years.
- Google, Yahoo!, Amazon, Microsoft, HP, Apple, and IBM are all companies have invested billions of dollars in datacentre construction and automation.
- Huge volumes of hardware, software, and database resources in these data centers can be allocated dynamically to millions of Internet users simultaneously.
- Guaranteed QoS and cost-effectiveness.

3.2 Virtualization

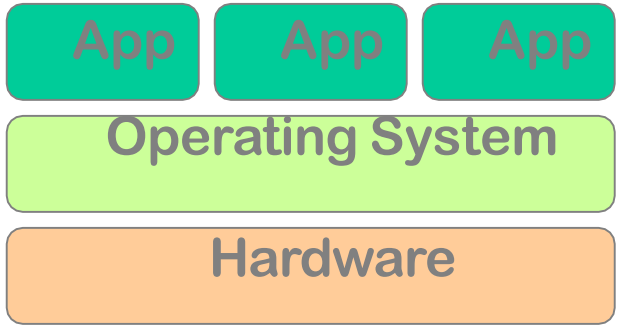

Traditional computing environment

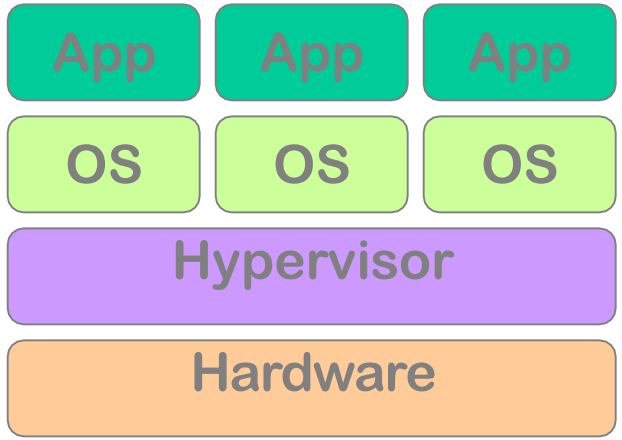

Virtualized computing environment

- **Hypervisor** manages multiple operating systems (or multiple instances of the same operating system) on a single computer system.
- The hypervisor manages the system's processor, memory, and other resources to allocate what each operating system requires.
- Hypervisors are designed for a particular processor architecture and may also be called **virtualization managers**.

A hypervisor or virtual machine monitor (VMM) is a piece of computer software, firmware or hardware that creates and runs virtual machines.

Benefits of Virtualization

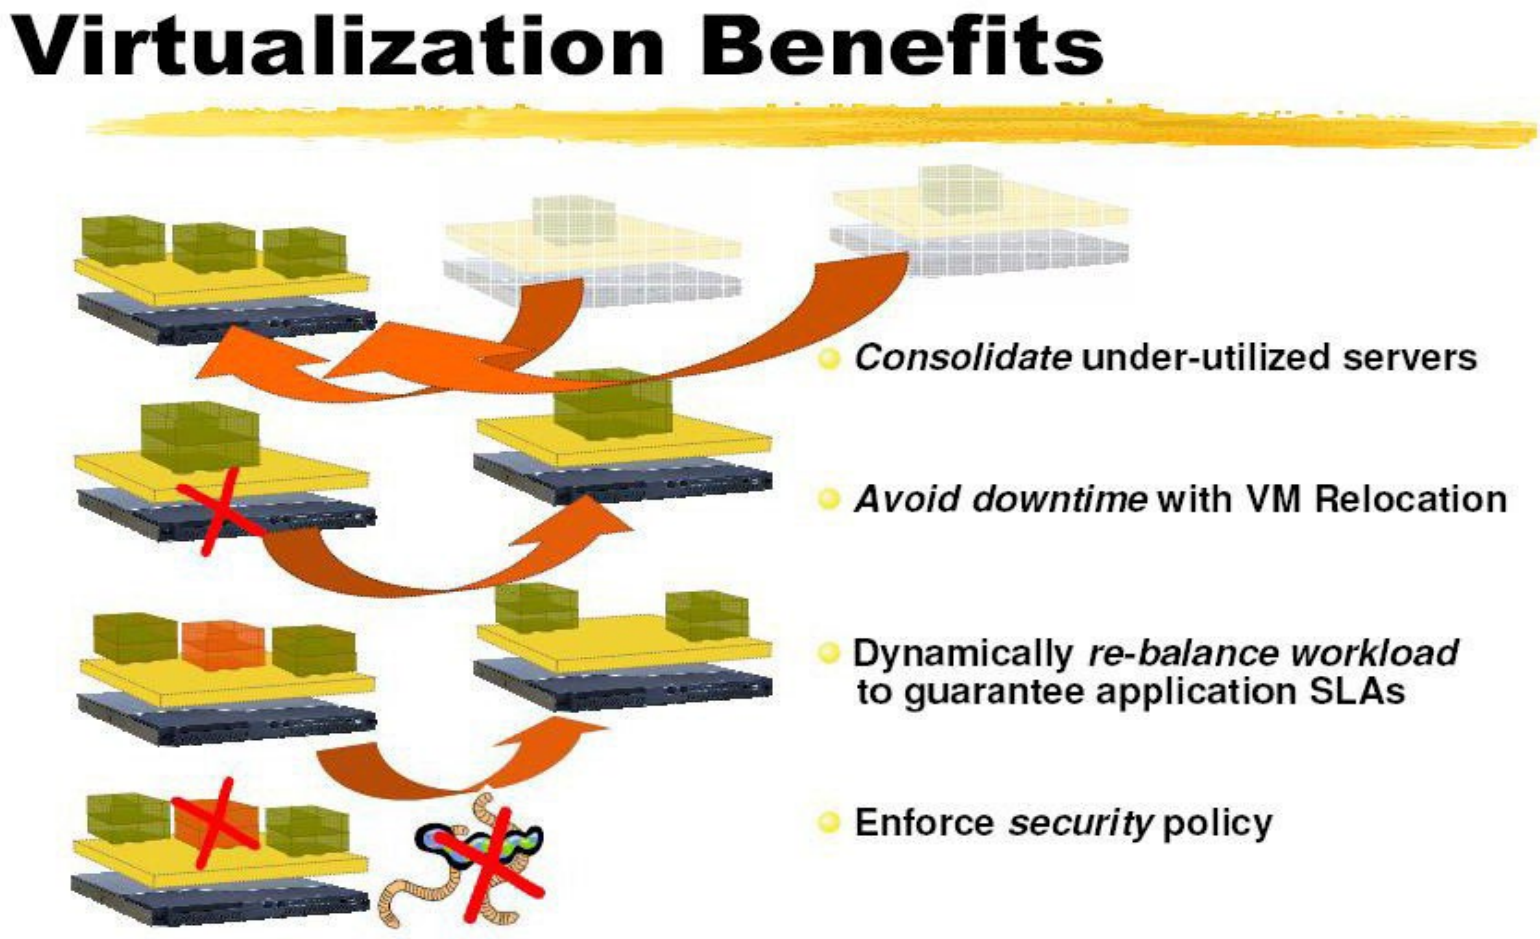

Benefits of Virtualization & Cloud Economics

Pay by use instead of provisioning for peak

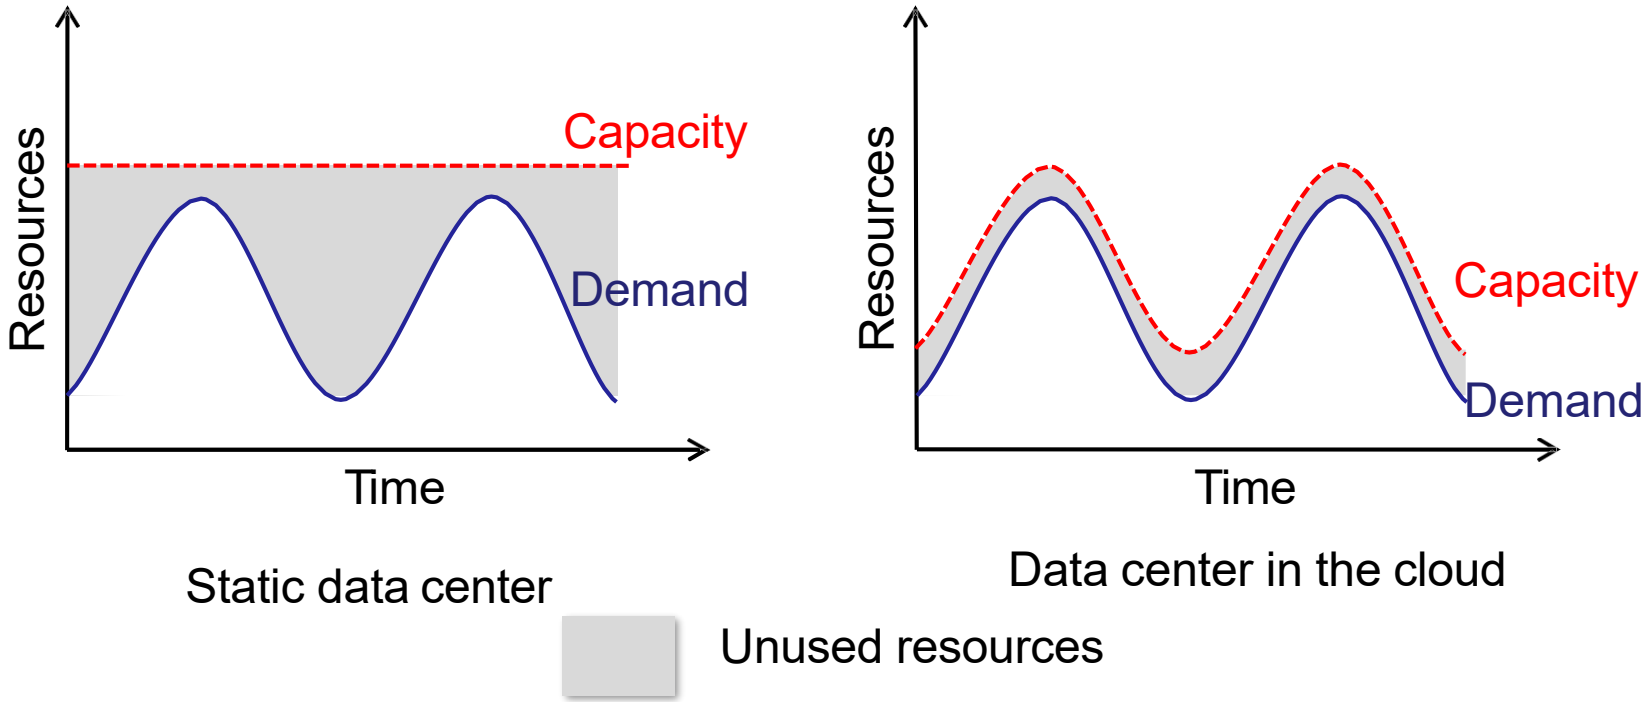

Benefits of Virtualization & Cloud Economics

Heavy penalty for under-provisioning

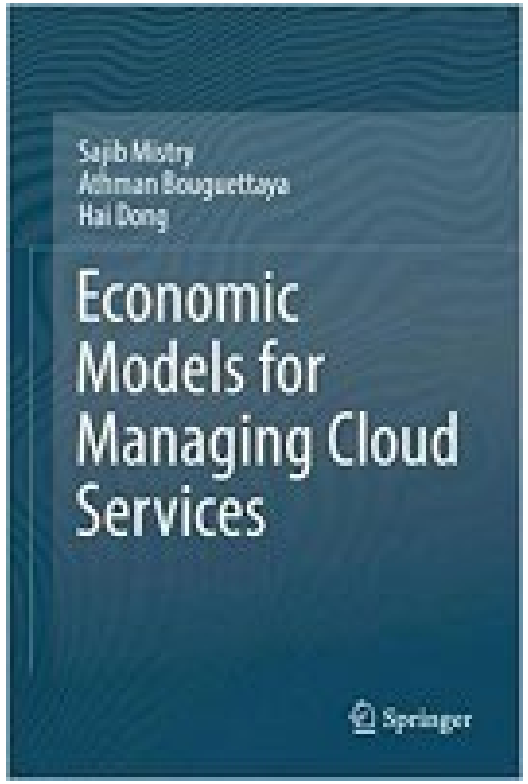

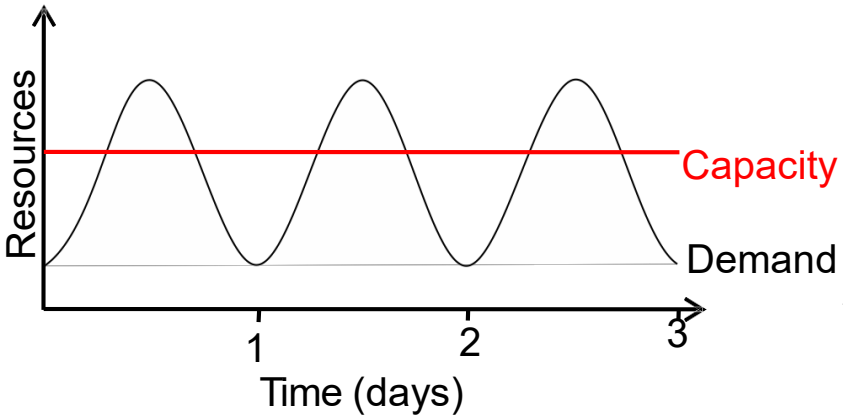

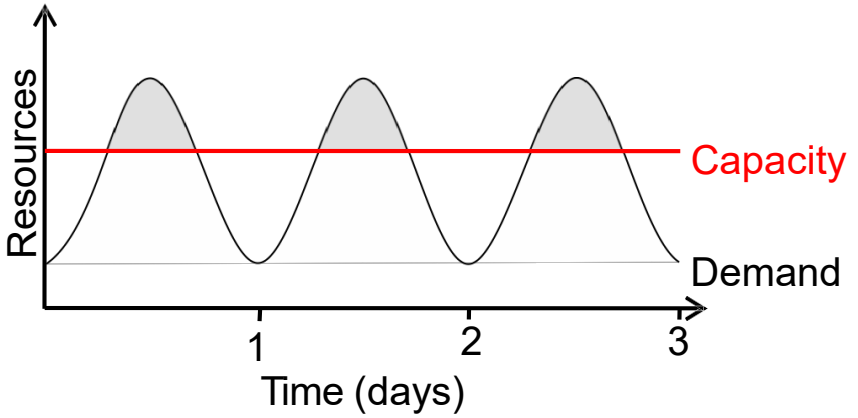

Lost revenue

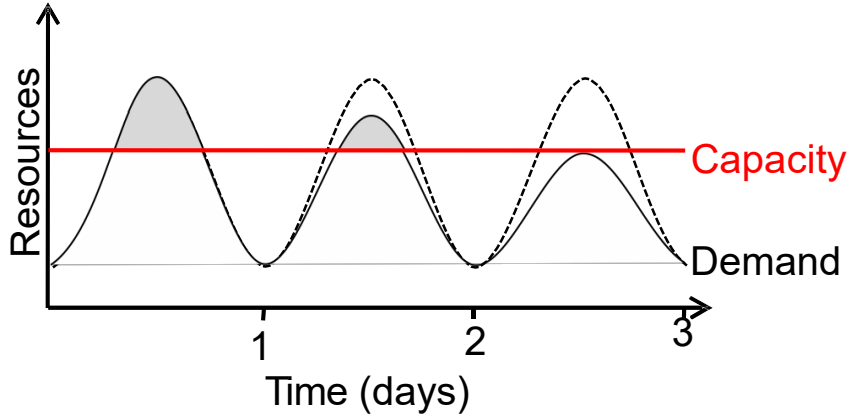

Lost users

https://www.springer.com/gp/book/9783319738758

3.3 Pay as you go

We calculate users' expenses when they execute a task in Amazon:

Pay as you Go

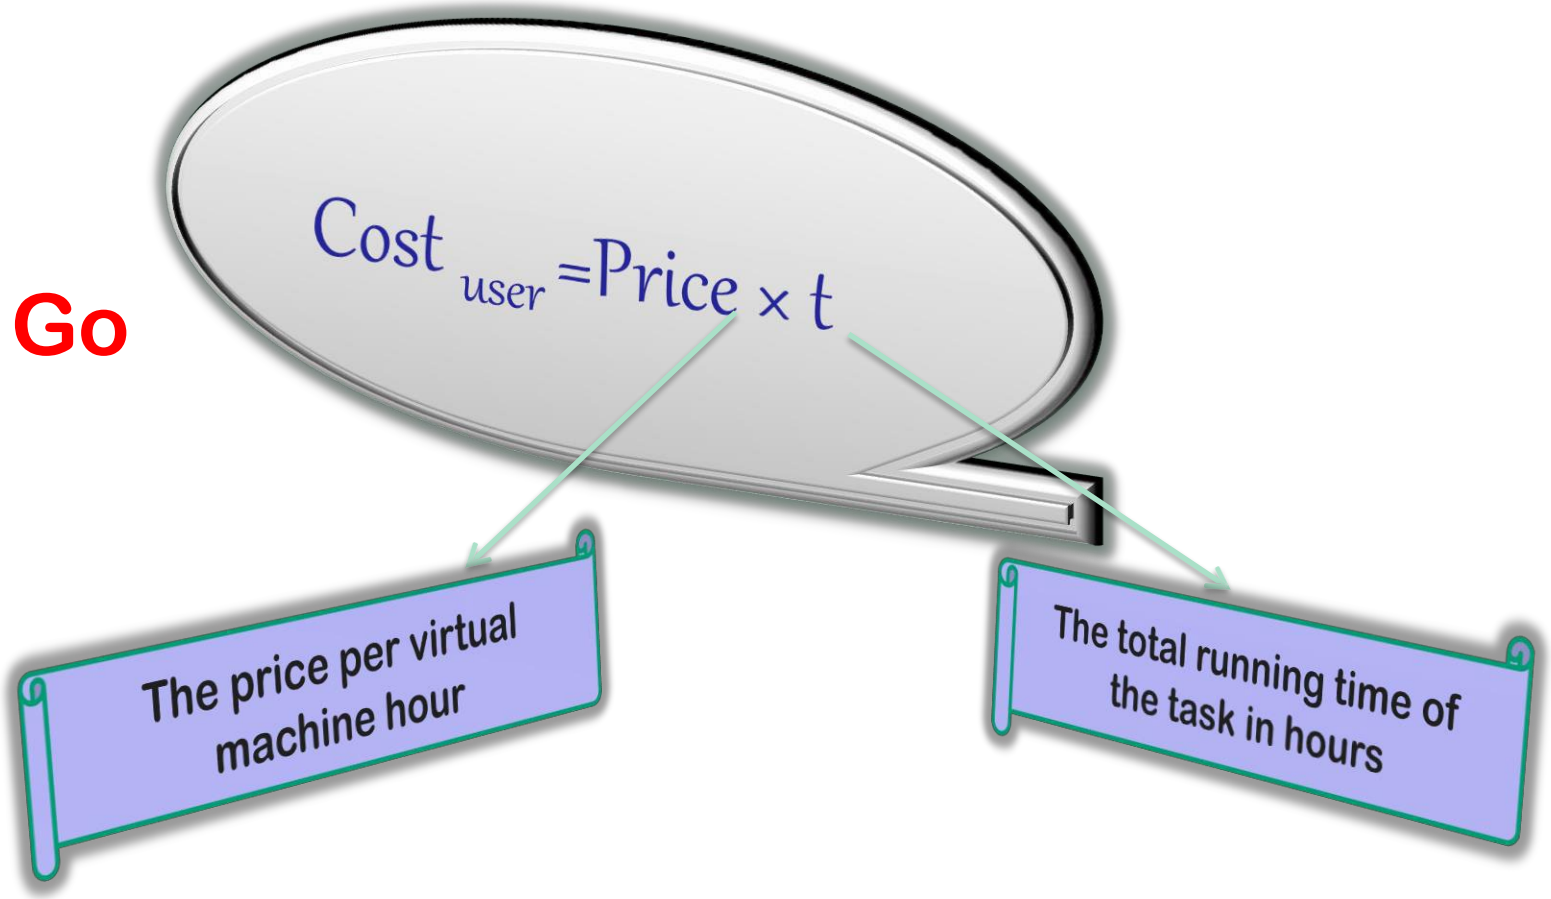

3.4 On Demand Self-Service

- A consumer can unilaterally provision computing capabilities, such as server time and network storage, as needed automatically without requiring human interaction with each service provider.
- Example: AWS Management Console

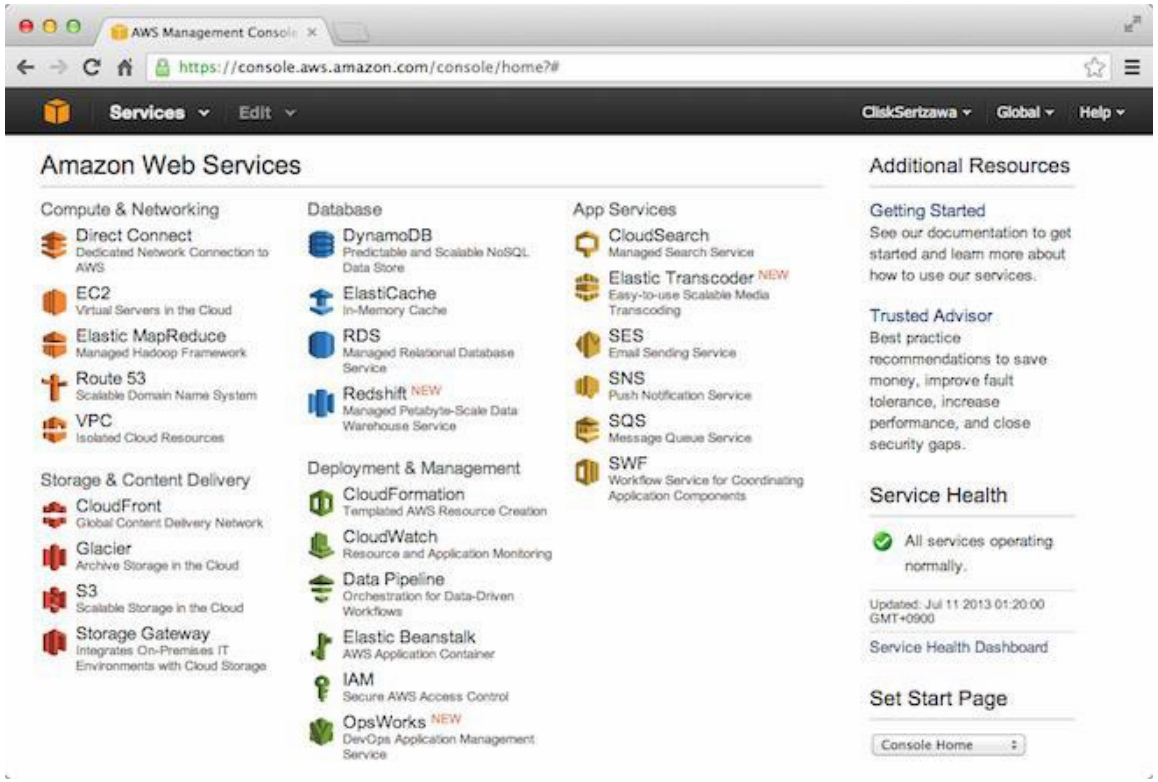

3.5 Broad Network Access

- Capabilities are available over the network and accessed through standard mechanisms that promote use by heterogeneous thin or thick client platforms.
- Example: Mobile phones, tablets, laptops, computers

3.6 Resource Pooling

- The provider's computing resources are pooled to serve multiple consumers using a multi-tenant model, with different physical and virtual resources dynamically assigned and reassigned according to consumer demand
- Examples: AWS availability zones

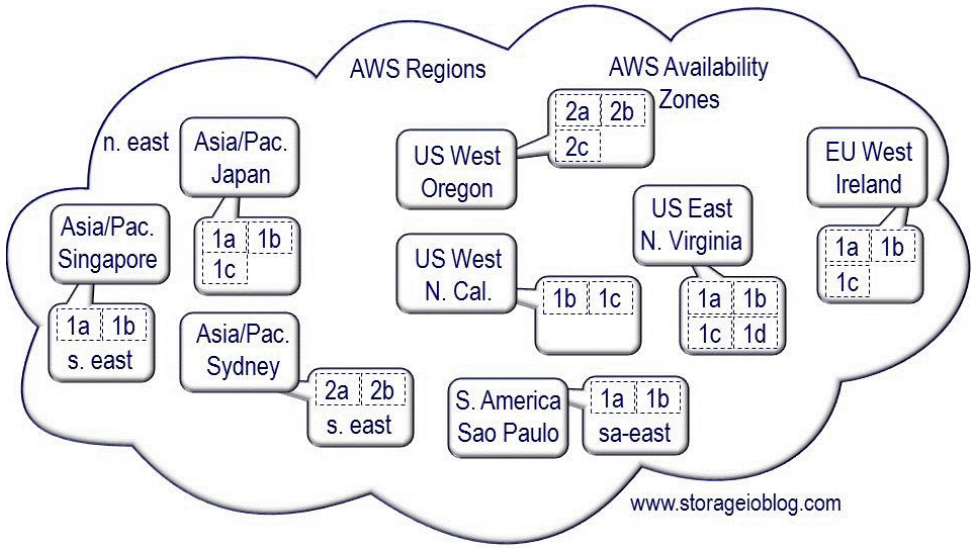

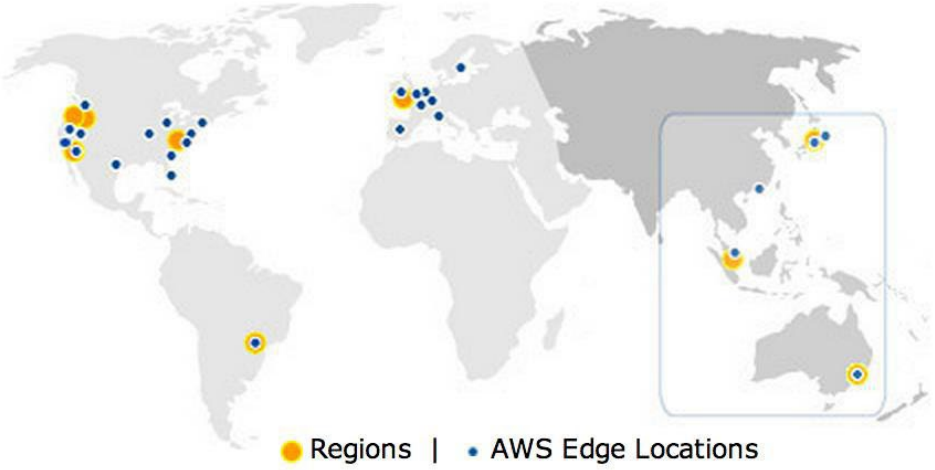

3.7 Rapid Elasticity

- Capabilities can be elastically provisioned and released, in some cases automatically, to scale rapidly outward and inward commensurate with demand
- Examples: AWS auto-scaling (will learn more in future lecture)

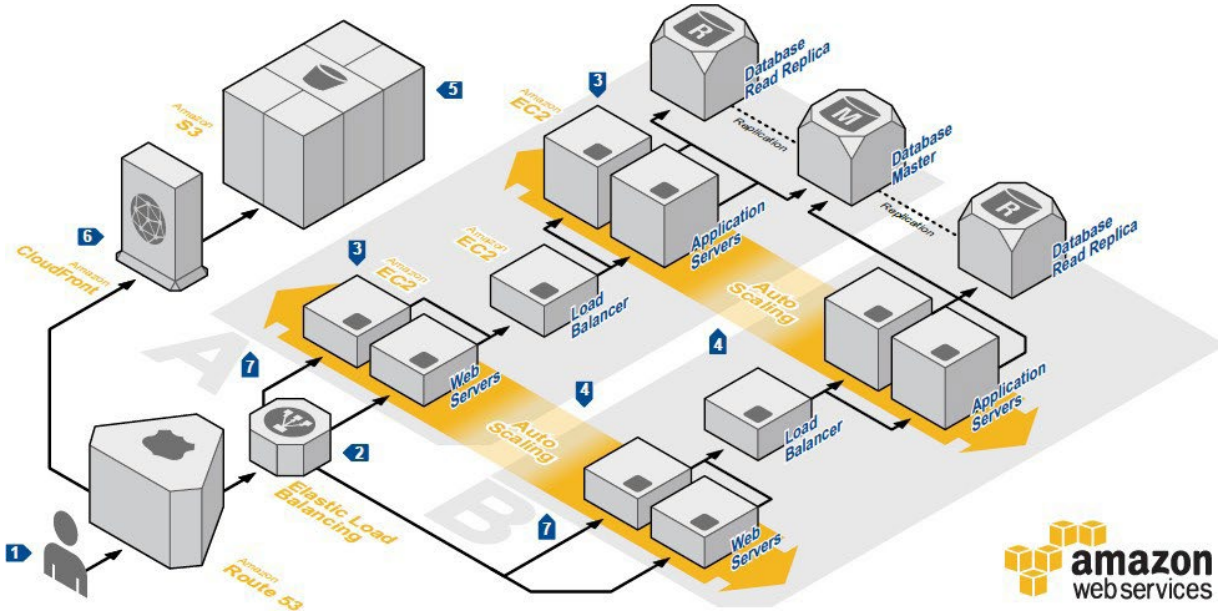

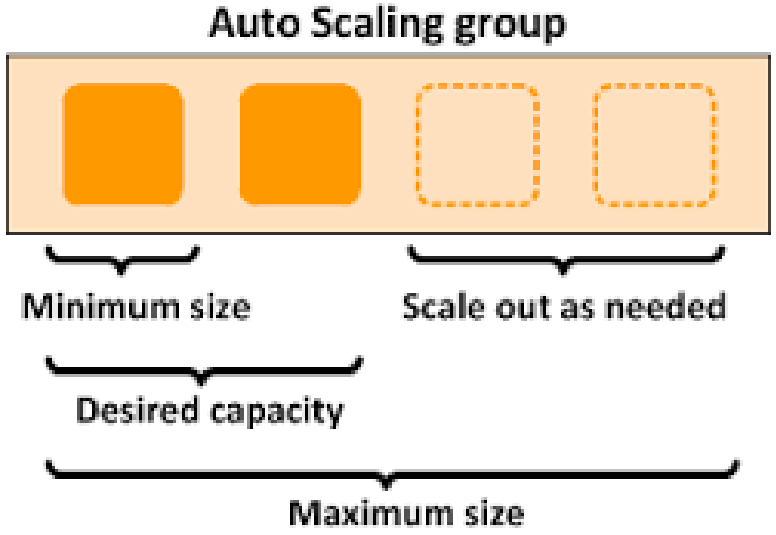

3.8 Measured Service

- Cloud systems automatically control and optimize resource by leveraging a metering capability
- Resource usage can be monitored, controlled, and reported, providing transparency for both the provider and consumer of the utilized service.
- Example: Pay-as-you-Go model

4 Cloud Service Models: IaaS, PaaS, SaaS

- Cloud Infrastructure as a Service (IaaS)
  - Rent processing, storage, network capacity, and other fundamental computing resources
- Cloud Platform as a Service (PaaS)
  - Deploy customer-created applications to a cloud
- Cloud Software as a Service (SaaS)
  - Use provider's applications over a network
- To be considered "cloud" they must be deployed on top of cloud infrastructure that has the key characteristics

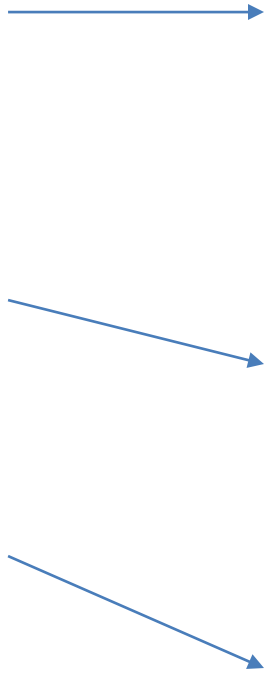

- Utility computing
  - Why buy machines when you can rent cycles?
  - Examples: Amazon's EC2, GoGrid, AppNexus
- Platform as a Service (PaaS)
  - Give me nice API and take care of the implementation
  - Example: Google App Engine
- Software as a Service (SaaS)
  - Just run it for me!
  - Example: Gmail, Facebook

Cloud Service Model Hierarchy

SalesForce CRM

LotusLive

Google App

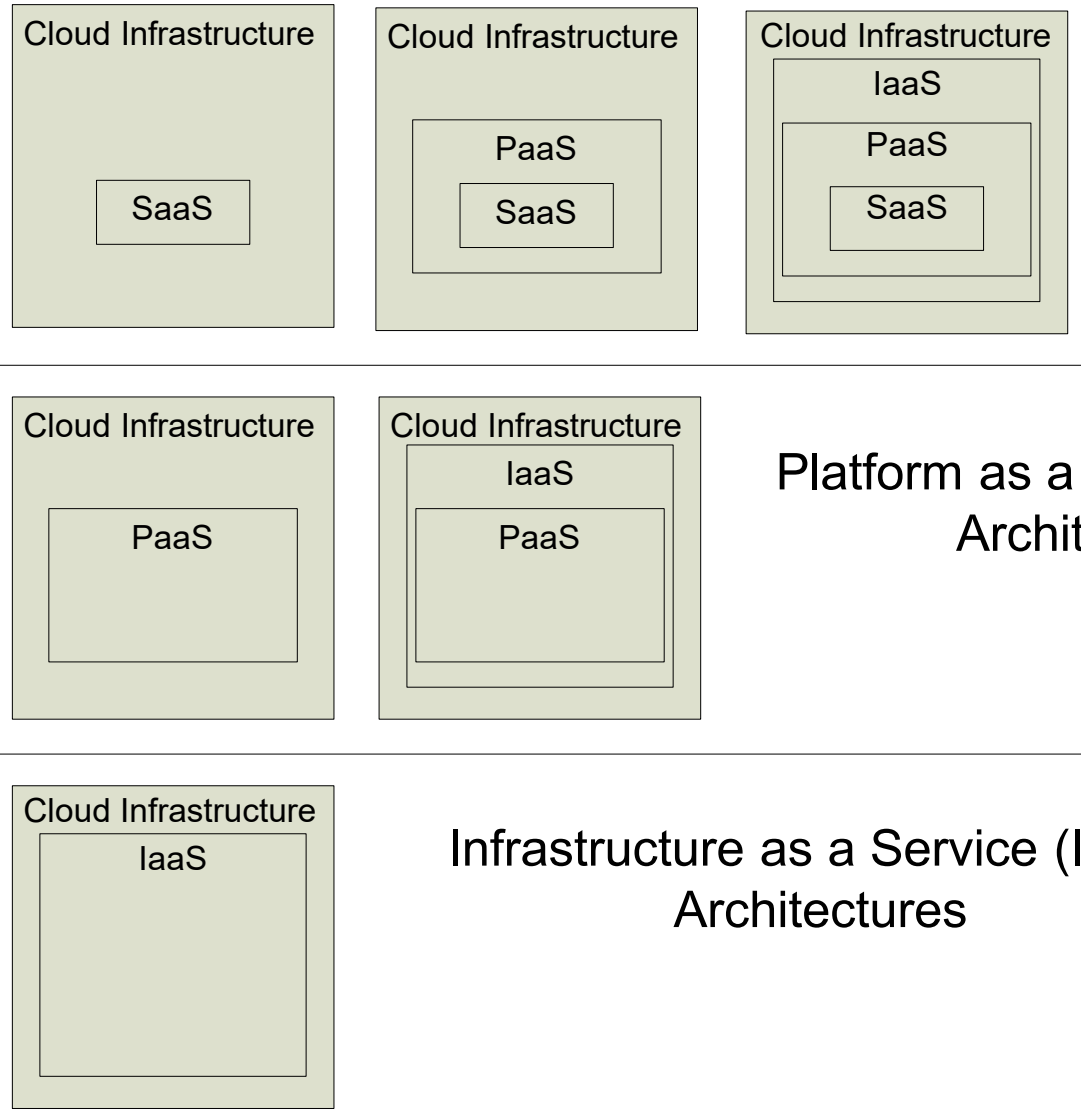

Software as a Service (SaaS) Architectures

Platform as a Service (PaaS) Architectures

Infrastructure as a Service (IaaS) Architectures

5 Cloud Deployment Models

- **Private cloud**
  - enterprise owned or leased. Cloud infrastructure for single org only, may be managed by the org or a 3rd party, on- or off-premise
- **Community cloud**
  - shared infrastructure for specific community. Shared by several orgs that have shared concerns, managed by org or 3rd party
- **Public cloud**
  - Sold to the public, mega-scale infrastructure. Cloud infrastructure is available to the general public, owned by org selling cloud services
- **Hybrid cloud**
  - composition of two or more cloud deployment models

Cloud Orchestration

- Managing and controlling all your cloud environments

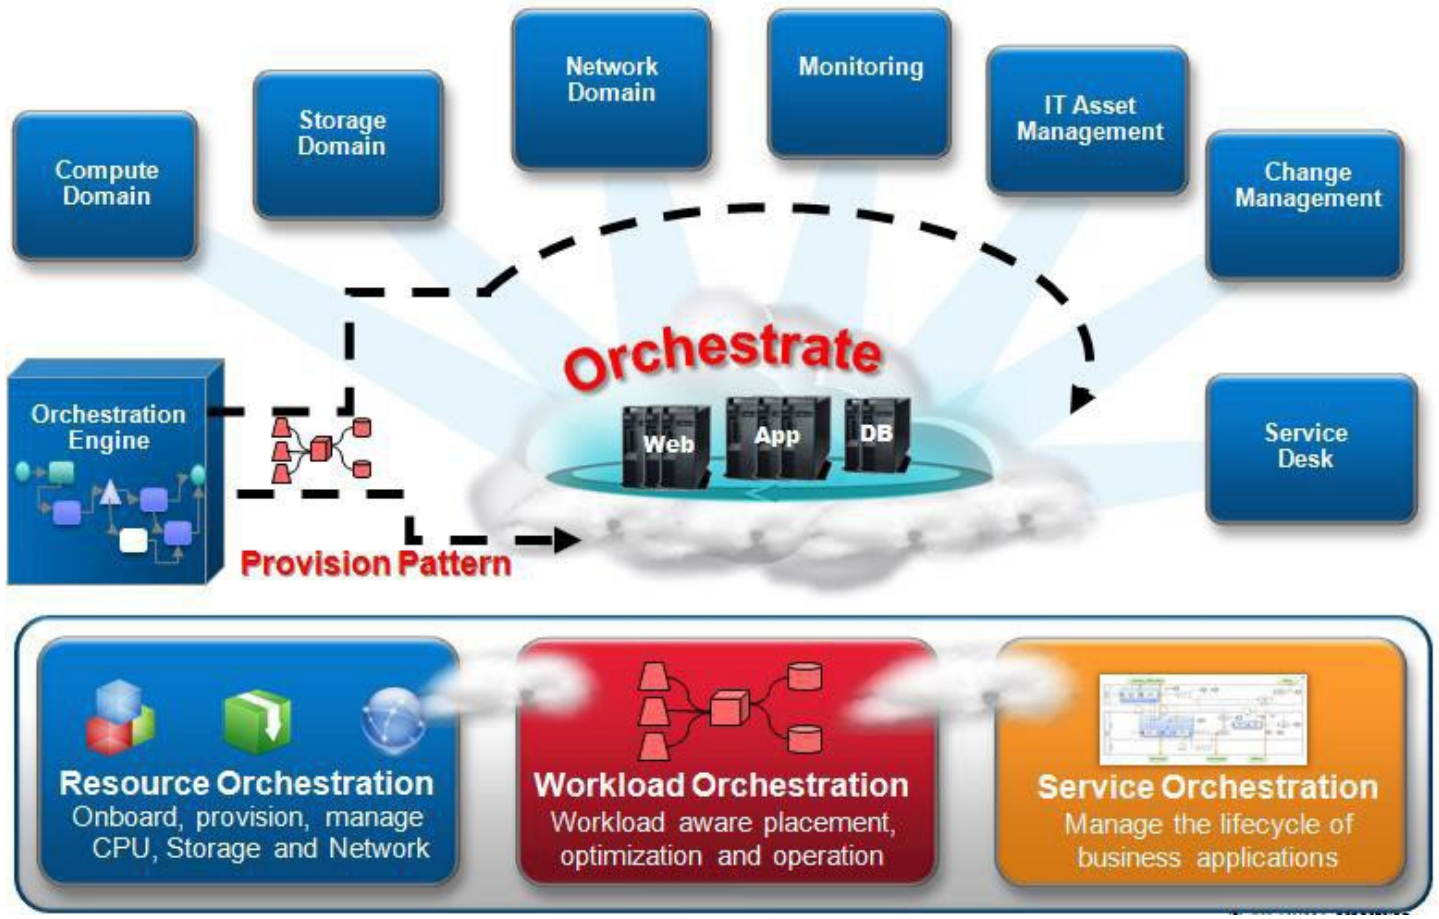

Orchestrated Cloud Environments